Inference from an L1 and L2 Perspective

Author: Jeremy H. Carroll **Date:** March 2026 **Version:** v1.0 (Pre-Print)

Abstract

In the foundational ℓ^1 verification framework, inference is fundamentally redefined. It is not solely the probabilistic prediction of the next token, as seen in standard ℓ^2 large language models. Rather, it can be modeled as a **morphism** in a category of reasoning states. This paper defines the precise topological mechanics of autonomous inference, contrasting the continuous ℓ^2 generation of token probabilities against the discrete ℓ^1 embedding of causal reasoning steps into a globally commutating directed acyclic graph.

---

1. Inference as a Morphism

An individual step of inference acts as a directed arrow that transforms one `ReasoningState` (representing the source context, domain, and constraint boundary) into a new `ReasoningState` (the target).

Every inference step is exclusively executed by a specific Lambda operator and carries a defined `MorphismKind` (such as a computation, generation, or verification step). By modeling reasoning categorially, every operational move an AI takes is mathematically typed, tracked, and structurally composable with subsequent or preceding steps.

---

2. Topological Validation via Commutativity

Because inference is treated as a sequence of formal morphisms, its validity is tested geometrically using **commutative diagrams**.

If the system applies two different chains of reasoning to navigate from concept A to concept C—for example, taking a direct inductive path h versus a two-step deductive path $g \circ f$—the results must perfectly match:

$$(g \circ f)(a) \equiv h(a)$$

If the diagram fails to commute, it acts as a topological tear detector, proving that two distinct parts of the system mathematically disagree about objective reality. In an ℓ^2 paradigm, non-commuting logic is statistically averaged out. In the ℓ^1 paradigm, the failure to commute immediately spikes the coboundary defect, halting the assumption that the logic is coherent.

2.1 Theorem (Non-Commutativity Implies ℓ^1 Defect Lower Bound)

Let $\mathcal{P}$ be a causal poset, $s \in C^0(\mathcal{P}, F)$ a section, $\delta = d_0 s$ the ℓ^1 coboundary, and $I(s) = \|\delta\|_1$. Consider a commutative diagram candidate $A \xrightarrow{f} B \xrightarrow{g} C$ and $A \xrightarrow{h} C$. Define the **commutativity defect**:

$$\Delta(a) := (g \circ f)(a) - h(a)$$

Claim: If the diagram does not commute ($\Delta(a) \neq 0$), then $I(s) \geq \|\Delta(a)\|_1$.

Proof: Each morphism corresponds to an edge path in the poset ($p_1 : A \to B \to C$ and $p_2 : A \to C$). The coboundary accumulates along paths such that $\sum_{e \in p_1} \delta_e(s) = (g \circ f)(a)$ and $\sum_{e \in p_2} \delta_e(s) = h(a)$. Subtracting these yields a cycle:

$$\sum_{e \in p_1} \delta_e(s) - \sum_{e \in p_2} \delta_e(s) = \Delta(a)$$

By the triangle inequality:

$$\|\Delta(a)\|_1 \leq \sum_{e \in p_1 \cup p_2} \|\delta_e(s)\|_1 \leq I(s)$$

Conclusion: Non-commutativity forces a strictly positive ℓ^1 defect lower bound. Commutativity equates to zero defect locally; non-commutativity ensures a structurally irreducible defect cost.

---

3. The Exact Mathematical Definition of a Hallucination

When an inference step produces an unresolvable contradiction (a non-zero ℓ^1 coboundary defect), it fails to structurally embed into the global constraint structure of the domain graph.

In this categorical framework, **a hallucination is formally defined as an inference step that cannot validly embed into the causal poset**. It is not a statistical anomaly or "bad data"; it is a massive geometric breakdown. The generated edge physically violates the commutative geometry of the object-level reasoning space.

---

4. Branching and the `ReasoningGraph`

As inference progresses, it generates a `ReasoningGraph`—a directed acyclic graph (DAG) where nodes are reasoning artifacts (prompts, outputs, refinements) and edges encode the formal parent-child derivation.

When the system faces uncertainty and diverges to explore multiple inference paths simultaneously, this branching transforms the standard reasoning DAG into **a stratified branching structure with cohomological obstruction depth**, suggestive of derived stack constructions. In this mathematically layered space, the cohomological depth (the H^1 measurement) explicitly tracks how many of these branches are structurally obstructed from gluing back together into a single, globally consistent truth.

---

5. Resolution of Obstructed Inference

If an inference path hits a dead end (measuring an $H^1 > 0$ geometric obstruction), the ℓ^1 system does not average out the error to find a "phantom point." Instead, it natively resolves the conflict by relying on an explicit mechanism:

5.1 Autopoietic Control

If an objective obstruction is identified, the system triggers the Autopoietic Iterator. This metacognitive control mechanism inserts new structural constraints (a specific `ReasoningMove`) designed to bridge the unresolvable gap. This operation corresponds to graph augmentation steps that iteratively reduce the defect until systemic consistency is achieved.

5.2 Theorem (Autopoietic Resolution of Non-Commutativity)

Every non-commuting diagram in the modeled reasoning graph induces a positive ℓ^1 defect. Under Autopoietic dynamics, the iterator is mathematically forced to act on non-commuting diagrams to reduce the defect: - **Repairable (Case A):** There exists an expansion Λ such that $I(\mathcal{P} \cup \Lambda, s) < I(\mathcal{P}, s)$. The iterator inserts edges, forcing the diagram to become commutative. - **Topologically Protected (Case B):** No such Λ exists ($I^* = \Phi([\delta]) > 0$); non-commutativity persists and becomes a stable structural invariant.

5.3 Structural Expansion vs. Fixed-Architecture Optimization

Unlike fixed-architecture optimization systems (Standard ML, MCTS, RL) that optimize values or policies over a fixed state space and handle contradictions statistically, the Autopoietic Iterator performs descent over both state and structure. It resolves inconsistency by modifying the underlying causal graph (pushout) rather than averaging over it.

---

6. Monotone Dynamics of Inference

In this topological framework, the ℓ^1 coboundary norm defines a monotone quantity analogous to entropy gradients in physical systems. The ℓ^1 coboundary norm

formally measures topological tension across the `Reasoning Graph`, ensuring inference is mathematically constrained by monotonic descent principles.

The Omega Engine executes this explicitly across four internal modules:

6.1 Epistemic Tension as Sheaf Restrictions

Abstract inference states are modeled as local sections of a pre-sheaf over the inference phase space. The ℓ^1 coboundary norm calculates the absolute mismatch across each overlapping edge in the complex. * **The Calculation:** $I(s) = \sum |\delta_e(s)|$. * **Interpretation:** This ℓ^1 tension measures structural disequilibrium—the exact degree to which local claims fail to glue into a globally coherent macrostate.

6.2 The Cohomological Obstruction (`stat_mech.rs`)

When the system projects a highly complex branching state into a single verifiable inference path, it evaluates the exact information differential. * **The Calculation:** H^1 obstruction scaling. * **Interpretation:** This evaluates the "information cost" of a reasoning projection. The ℓ^1 obstruction here is mathematically guaranteed to be ≥ 0, representing the irreducible tension that prevents the AI from collapsing contradictory logics into a localized point.

6.3 Inference Trace Evaluation (`compute_cognitive_z.rs`)

The engine can evaluate individual epistemic traces by treating them as weighted paths in the reasoning DAG. The "action" equivalent $S[s]$ of a specific inference sequence is defined rigorously as its total ℓ^1 topological defect: $S[s] = \|\delta s\|_1$. The logged `perception.entropy` metric serves as a heuristic evaluation proxy for tracking this coboundary cost during generation.

6.4 Coupled within the Lyapunov Control Function (`crf_tensor.rs`)

The system bounds active generation using a Lyapunov-analogous descent function that restricts generation to monotonic pathways. * **The Calculation:** $V(s) \propto I(s) + \gamma(\text{divergence bounds})$. * **Interpretation:** The control function bounds the maximum random topological diffusion permitted by the system. If divergence spirals ($V > 0$), topological negative feedback forces compression along the constraint graph.

Hallucination in structural intelligence is a **cohomological obstruction**. The inference steps provide the 0-cochain values (the nodes), and the ℓ^1 norm mathematically measures the consistency limits across the causal edges.

---

7. Calculating Inference: ℓ^1 vs. ℓ^2

The core difference between calculating inference in ℓ^1 geometry versus standard ℓ^2 geometry comes down to **sparsity vs. diffusion**.

- **ℓ^1 Inference (Sparsity & Isolation):** In ℓ^1 geometry, inference errors are evaluated by computing the exact coboundary defect $I(s) = \sum \|\delta_e(s)\|_1$, which sums the absolute discrepancies across all reasoning edges. When resolving errors, ℓ^1 calculation finds the **median** rather than the mean, ensuring that the system does not interpolate between contradictory claims. It encourages the sparsity of data and forces unresolved contradictions to concentrate onto lower-dimensional boundaries (min-cuts), structurally localizing inconsistency.
- **ℓ^2 Inference (Diffusion & Smoothing):** ℓ^2 inference evaluates errors using the sum of squares, inherently allowing contradictions of opposite signs to partially cancel each other out. In the Rust engine, ℓ^2 optimization behaves like a graph Laplacian or a heat equation; it diffuses the residual error globally across the entire network, smoothing out gradients but mathematically permitting the diffusion of inconsistency across structural boundaries.
- **Calculating the Difference (The Tension):** The mathematical difference between the two is explicitly calculated using the **Inflation Ratio**. The system evaluates the raw defect using an ℓ^2 Hodge decomposition (splitting it into exact and harmonic components via orthogonal projection) and compares the inflated cost against the true ℓ^1 quotient norm Φ. Mapping a sparse exact defect into an ℓ^2 harmonic space universally inflates the mathematical cost, strictly approaching a ratio of 3 as observed in cyclic graph constructions (see Paper 007b).

---

8. Mathematical Tools in the Rust Codebase

To perform these evaluations efficiently, the `omega_core` codebase utilizes several rigorous mathematical and topological engines:

8.1 The `CoboundaryNorm` Engine (`l028`)

This is the foundational tool for ℓ^1 inference. It takes a `FinitePoset` and a `Section` (the values assigned to the nodes) and strictly computes `compute_l1_rigidity()`, natively generating the total absolute defect across all covering edges.

8.2 ℓ^1-Native Hodge Decomposition & ISTA (`l302`)

Instead of relying on the classical ℓ^2 Hodge projection—which uses Euclidean inner products—the codebase implements `L1NativeDecomposition`. To compute the primal minimizer ($\min_f \|\delta - d_0 f\|_1$), it uses `solve_primal_ista`, an

Iterative Soft-Thresholding Algorithm that natively handles ℓ^1 projection to gracefully isolate exact gauge artifacts from irreducible topological tears.

8.3 The Autopoietic Iterator / Wasserstein Descent (`l166`)

To calculate the structural evolution of inference, this tool implements discrete Wasserstein-1 gradient descent. Instead of global ℓ^2 smoothing, it iteratively detects the worst ℓ^1 tear, calculates local ℓ^1-optimal medians on star neighborhoods, and surgically augments the poset with new geometric edges.

8.4 Linear Programming (LP) and Network Simplex (`l097`)

To definitively calculate the irreducible domain complexity floor (Φ), the codebase frames the ℓ^1 quotient norm as a Linear Program (LP). The LP solver constructs the exact primal-dual pairing by evaluating bounded divergence-free circulations against the defect field, calculating the ultimate topological obstruction in polynomial time $O(|E|^{1.5})$.

8.5 Spectral Graph Laplacian (`l097`)

When specifically measuring ℓ^2 geometries (or evaluating precisely why they structurally fail), the system uses the `GraphLaplacian` ($L = D - A$) to calculate the Dirichlet energy $E(\rho) = Tr(\rho L)$. It applies `spectral_diffusion` (the heat kernel e^{-tL}) to simulate how ℓ^2 mechanisms smoothly spread probability states across graph neighborhoods. It also utilizes the `lanczos` algorithm to efficiently extract top-k eigenvalues for extremely fast von Neumann entropy approximations.

---

9. Conclusion

Inference in the verification framework is the continuous, dynamic mapping of morphisms across a causal poset, tightly governed by the ℓ^1 coboundary norm to ensure every logical step mathematically aligns with objective thermodynamic reality. This demonstrates structural limitations of ℓ^2 autoregressive inference in the presence of contradictory constraints, motivating ℓ^1-based alternatives.